

REVIEW of PAPER # 17 BY REVIEWER # 10

Fill in each _____ with a number on a scale of 1-5.

(1=Poor) (2=Fair) (3=Acceptable) (4=Good) (5=Excellent)

1. Presentation: 4

1.1. Organization: 4

1.2 Grammar and spelling: 4

2. Completeness (Strength of Content- Missing key items?): 3

3. Technical Correctness: 3

4. Proper Referencing: 1

5. "Coolness" / Originality: 3

6. Comments to Author (Suggestions for Improvement):

A few items:

- Embed citations into the text to substantiate claims
- Catch misspellings ('detrunents', 'referred', 'indenticle')
- Abstract should include main conclusions

Overall, it was an interesting paper to read. I liked the Xen discussion and some of the additional insights included in that section. It was also interesting to see the various performance concerns between the different systems. However, there were some correctness errors. Specifically, I don't know if I would consider Xen, Exokernel, or Virtual Ghost operating systems. Also, I am not convinced your definition of paravirtualization is correct. Check the wikipedia definition for paravirtualization to see what I mean. There are a few other minor items that might have been cleared up if I had embedded citations to reference. The section on Virtual Ghost felt a little clunky. For instance, it was not clear to me how allowing applications to choose their own encryption algorithms was a performance benefit. The conclusion could have had a bit more substance and been a bit stronger.

